

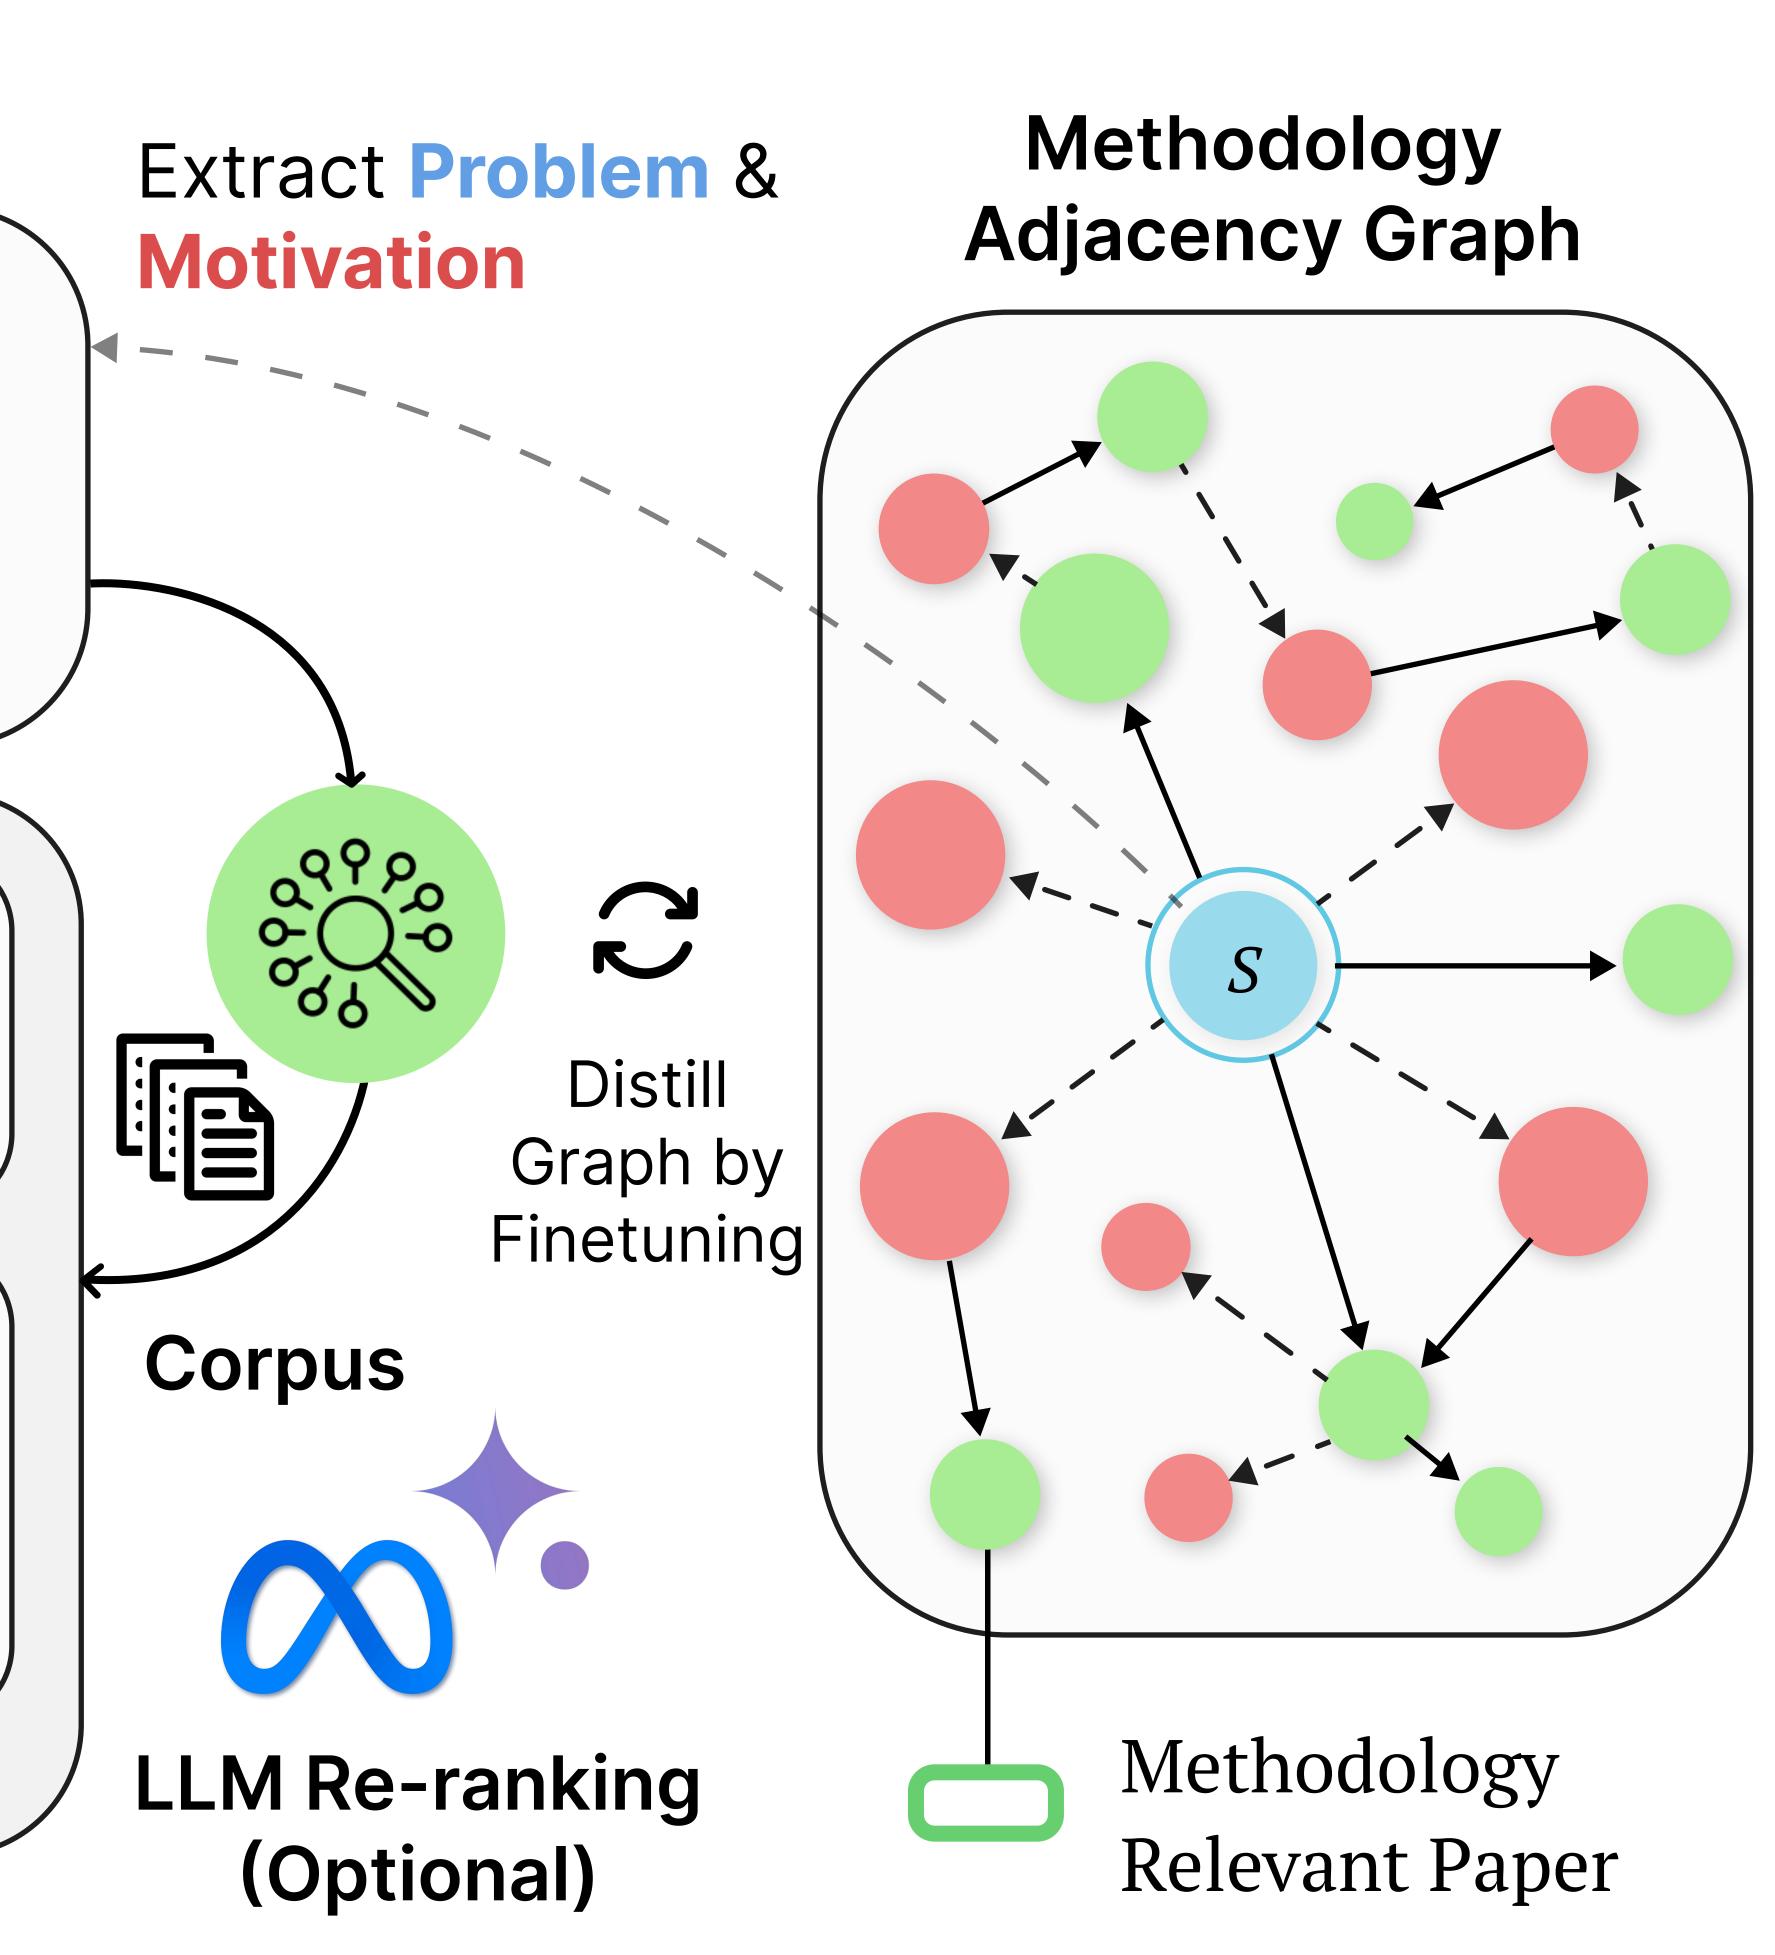
Proposal: The research aims to effectively compress BERT, a large language model, while maintaining its performance. The motivation is to improve the efficiency of BERT by reducing its size without sacrificing performance, improve existing knowledge distillation approaches by introducing deeper level of interaction between the original and compact models.

Title: DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter

Abstract: ...a method to pre-train a smaller general purpose language representation model ...introduce a triple loss combining language modeling, distillation and cosine-distance losses...

Title: Patient Knowledge Distillation for BERT Model Compression

Abstract: ...alleviate this resource hunger in large-scale ...high demand for computing resources in training such models hinders... ...patient distillation schemes enable the exploitation of rich information in the teacher's hidden layers...



Relevant Paper